

Redes neuronales artificiales con Keras

Este notebook contiene todo el código fuente requerido para la solución de los talleres propuestos.



[Ejecutar en Google Colab](#)

Configuración

Primero, importemos algunos módulos comunes, asegurémonos de que Matplotlib traza las figuras en línea y preparemos una función para guardar las figuras. También verifiquemos que Python 3.5 o posterior esté instalado (aunque Python 2.x puede funcionar, está obsoleto, por lo que recomendamos encarecidamente que use Python 3 en su lugar), así como Scikit-Learn ≥ 0.20 y TensorFlow ≥ 2.0 .

In [1]:

```
# Python  $\geq 3.5$  es requerido
import sys
assert sys.version_info >= (3, 5)

# Scikit-Learn  $\geq 0.20$  es requerido
import sklearn
assert sklearn.__version__ >= "0.20"

try:
    # %tensorflow_version solo existe en Colab.
    %tensorflow_version 2.x
except Exception:
    pass

# TensorFlow  $\geq 2.0$  es requerido
import tensorflow as tf
assert tf.__version__ >= "2.0"

# Importar librerías comunes
import numpy as np
import os

# para que la salida de este notebook sea estable en todas las ejecuciones
np.random.seed(42)

# Para dibujar figuras estéticas
%matplotlib inline
import matplotlib as mpl
import matplotlib.pyplot as plt
mpl.rc('axes', labelsiz=14)
mpl.rc('xtick', labelsiz=12)
mpl.rc('ytick', labelsiz=12)

# En donde almacenar las figuras
PROJECT_ROOT_DIR = "."
CHAPTER_ID = "ann"
IMAGES_PATH = os.path.join(PROJECT_ROOT_DIR, "images", CHAPTER_ID)
os.makedirs(IMAGES_PATH, exist_ok=True)

def save_fig(fig_id, tight_layout=True, fig_extension="png", resolution=300):
    path = os.path.join(IMAGES_PATH, fig_id + "." + fig_extension)
    print("Saving figure", fig_id)
    if tight_layout:
        plt.tight_layout()
    plt.savefig(path, format=fig_extension, dpi=resolution)

# Ignorar las advertencias inútiles (consulte el número 5998 de SciPy)
import warnings
warnings.filterwarnings(action="ignore", message="^internal gelsd")
```

Perceptrones

Nota: establecemos `max_iter` y `tol` explícitamente para evitar advertencias sobre el hecho de que su valor predeterminado cambiará en futuras versiones de Scikit-Learn.

In [2]:

```
import numpy as np
from sklearn.datasets import load_iris
from sklearn.linear_model import Perceptron

iris = load_iris()
X = iris.data[:, (2, 3)] # largo del pétalo, ancho del pétalo
y = (iris.target == 0).astype(np.int)

per_clf = Perceptron(max_iter=1000, tol=1e-3, random_state=42)
per_clf.fit(X, y)

y_pred = per_clf.predict([[2, 0.5]])
```

In [3]:

```
y_pred
```

Out[3]:

```
array([1])
```

In [101]:

```
a = -per_clf.coef_[0][0] / per_clf.coef_[0][1]
b = -per_clf.intercept_ / per_clf.coef_[0][1]

axes = [0, 5, 0, 2]

x0, x1 = np.meshgrid(
    np.linspace(axes[0], axes[1], 500).reshape(-1, 1),
    np.linspace(axes[2], axes[3], 200).reshape(-1, 1),
)
X_new = np.c_[x0.ravel(), x1.ravel()]
y_predict = per_clf.predict(X_new)
zz = y_predict.reshape(x0.shape)

plt.figure(figsize=(10, 4))
plt.plot(X[y==0, 0], X[y==0, 1], "bs", label="No es Iris-Setosa")
plt.plot(X[y==1, 0], X[y==1, 1], "yo", label="Iris-Setosa")


plt.plot([axes[0], axes[1]], [a * axes[0] + b, a * axes[1] + b], "k-", linewidth=3)
from matplotlib.colors import ListedColormap
custom_cmap = ListedColormap(['#9898ff', '#fafab0'])

plt.contourf(x0, x1, zz, cmap=custom_cmap)
plt.xlabel("Longitud Pétalo", fontsize=14)
plt.ylabel("Ancho Pétalo", fontsize=14)
plt.legend(loc="lower right", fontsize=14)
plt.axis(axes)

save_fig("perceptron_iris_plot")
plt.show()
```

Saving figure perceptron_iris_plot





In []:

```
# Funciones de Activación
```

In [5]:

```
def sigmoid(z):
    return 1 / (1 + np.exp(-z))

def relu(z):
    return np.maximum(0, z)

def derivative(f, z, eps=0.000001):
    return (f(z + eps) - f(z - eps)) / (2 * eps)
```

In [102]:

```
z = np.linspace(-5, 5, 200)


plt.figure(figsize=(11,4))

plt.subplot(121)
plt.plot(z, np.sign(z), "r-", linewidth=1, label="Step")
plt.plot(z, sigmoid(z), "g--", linewidth=2, label="Sigmoid")
plt.plot(z, np.tanh(z), "b-", linewidth=2, label="Tanh")
plt.plot(z, relu(z), "m-.", linewidth=2, label="ReLU")
plt.grid(True)
plt.legend(loc="center right", fontsize=14)
plt.title("Funciones de Activación", fontsize=14)
plt.axis([-5, 5, -1.2, 1.2])

plt.subplot(122)
plt.plot(z, derivative(np.sign, z), "r-", linewidth=1, label="Step")
plt.plot(0, 0, "ro", markersize=5)
plt.plot(0, 0, "rx", markersize=10)
plt.plot(z, derivative(sigmoid, z), "g--", linewidth=2, label="Sigmoid")
plt.plot(z, derivative(np.tanh, z), "b-", linewidth=2, label="Tanh")
plt.plot(z, derivative(relu, z), "m-.", linewidth=2, label="ReLU")
plt.grid(True)
# plt.legend(loc="center right", fontsize=14)
plt.title("Derivativas", fontsize=14)
plt.axis([-5, 5, -0.2, 1.2])

save_fig("activation_functions_plot")
plt.show()
```

Saving figure activation_functions_plot

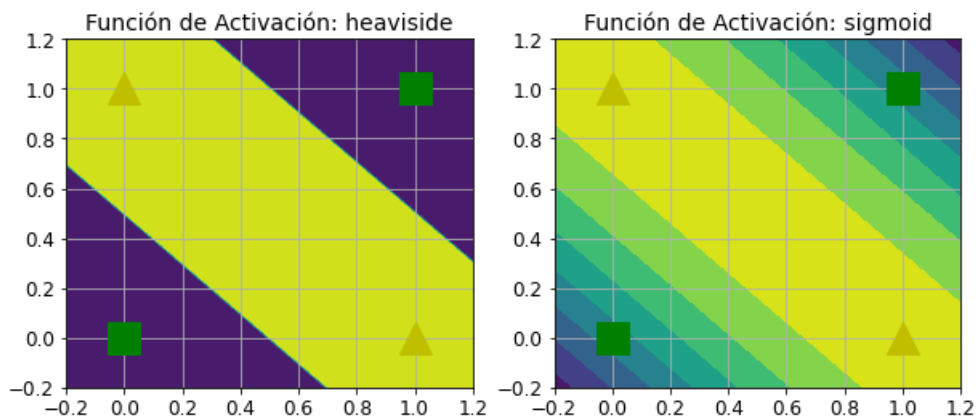


In [7]:

```
def heaviside(z):  
    return (z >= 0).astype(z.dtype)  
  
def mlp_xor(x1, x2, activation=heaviside):  
    return activation(-activation(x1 + x2 - 1.5) + activation(x1 + x2 - 0.5) - 0.5)
```

In [103]:

```
x1s = np.linspace(-0.2, 1.2, 100)  
x2s = np.linspace(-0.2, 1.2, 100)  
x1, x2 = np.meshgrid(x1s, x2s)  
  
z1 = mlp_xor(x1, x2, activation=heaviside)  
z2 = mlp_xor(x1, x2, activation=sigmoid)  
  
plt.figure(figsize=(10,4))  
  
plt.subplot(121)  
plt.contourf(x1, x2, z1)  
plt.plot([0, 1], [0, 1], "gs", markersize=20)  
plt.plot([0, 1], [1, 0], "y^", markersize=20)  
plt.title("Función de Activación: heaviside", fontsize=14)  
plt.grid(True)  
  
plt.subplot(122)  
plt.contourf(x1, x2, z2)  
plt.plot([0, 1], [0, 1], "gs", markersize=20)  
plt.plot([0, 1], [1, 0], "y^", markersize=20)  
plt.title("Función de Activación: sigmoid", fontsize=14)  
plt.grid(True)
```



Construyendo un clasificador de Imágenes

Primero, importemos TensorFlow y Keras.

In [9]:

```
import tensorflow as tf  
from tensorflow import keras
```

In [10]:

```
tf.__version__
```

Out[10]:

```
'2.3.1'
```

In [11]:

```
""" [11] :
```

```
keras.__version__
```

```
Out[11]:
```

```
'2.4.0'
```

Comencemos cargando el conjunto de datos de moda MNIST. Keras tiene una serie de funciones para cargar conjuntos de datos populares en `keras.datasets`. El conjunto de datos ya está dividido entre un conjunto de entrenamiento y un conjunto de prueba, pero puede ser útil dividir el conjunto de entrenamiento más para tener un conjunto de validación:

```
In [12]:
```

```
fashion_mnist = keras.datasets.fashion_mnist  
(X_train_full, y_train_full), (X_test, y_test) = fashion_mnist.load_data()
```

El conjunto de entrenamiento contiene 60.000 imágenes en escala de grises, cada una de 28x28 píxeles:

```
In [13]:
```

```
X_train_full.shape
```

```
Out[13]:
```

```
(60000, 28, 28)
```

Cada intensidad de píxel se representa como un byte (de 0 a 255):

```
In [14]:
```

```
X_train_full.dtype
```

```
Out[14]:
```

```
dtype('uint8')
```

Dividamos el conjunto de entrenamiento completo en un conjunto de validación y un conjunto de entrenamiento (más pequeño). También escalamos las intensidades de píxeles hasta el rango 0-1 y las convertimos en flotantes, dividiéndolas por 255.

```
In [15]:
```

```
X_valid, X_train = X_train_full[:5000] / 255., X_train_full[5000:] / 255.  
y_valid, y_train = y_train_full[:5000], y_train_full[5000:]  
X_test = X_test / 255.
```

Puede trazar una imagen usando la función `imshow()` de Matplotlib, con un `'binario'` mapa de color:

```
In [16]:
```

```
plt.imshow(X_train[0], cmap="binary")  
plt.axis('off')  
plt.show()
```



Las etiquetas son los ID de clase (representados como uint8), de 0 a 9:

```
In [17]:
```

```
y_train
```

```
Out[17]:
```

```
array([4, 0, 7, ..., 3, 0, 5], dtype=uint8)
```

Aquí están los nombres de clase correspondientes:

```
In [18]:
```

```
class_names = ["T-shirt/top", "Trouser", "Pullover", "Dress", "Coat",  
               "Sandal", "Shirt", "Sneaker", "Bag", "Ankle boot"]
```

Entonces, la primera imagen del conjunto de entrenamiento es un abrigo:

```
In [19]:
```

```
class_names[y_train[0]]
```

```
Out[19]:
```

```
'Coat'
```

El conjunto de validación contiene 5000 imágenes y el conjunto de prueba contiene 10000 imágenes:

```
In [20]:
```

```
X_valid.shape
```

```
Out[20]:
```

```
(5000, 28, 28)
```

```
In [21]:
```

```
X_test.shape
```

```
Out[21]:
```

```
(10000, 28, 28)
```

Echemos un vistazo a una muestra de las imágenes en el conjunto de datos:

```
In [22]:
```

```
n_rows = 4  
n_cols = 10  
plt.figure(figsize=(n_cols * 1.2, n_rows * 1.2))  
for row in range(n_rows):  
    for col in range(n_cols):  
        index = n_cols * row + col  
        plt.subplot(n_rows, n_cols, index + 1)  
        plt.imshow(X_train[index], cmap="binary", interpolation="nearest")  
        plt.axis('off')  
        plt.title(class_names[y_train[index]], fontsize=12)  
plt.subplots_adjust(wspace=0.2, hspace=0.5)  
save_fig('fashion_mnist_plot', tight_layout=False)  
plt.show()
```

Saving figure fashion_mnist_plot



In [23]:

```
model = keras.models.Sequential()
model.add(keras.layers.Flatten(input_shape=[28, 28]))
model.add(keras.layers.Dense(300, activation="relu"))
model.add(keras.layers.Dense(100, activation="relu"))
model.add(keras.layers.Dense(10, activation="softmax"))
```

In [24]:

```
keras.backend.clear_session()
np.random.seed(42)
tf.random.set_seed(42)
```

In [25]:

```
model = keras.models.Sequential([
    keras.layers.Flatten(input_shape=[28, 28]),
    keras.layers.Dense(300, activation="relu"),
    keras.layers.Dense(100, activation="relu"),
    keras.layers.Dense(10, activation="softmax")
])
```

In [26]:

```
model.layers
```

Out[26]:

```
[<tensorflow.python.keras.layers.core.Flatten at 0x135da4c0>,
<tensorflow.python.keras.layers.core.Dense at 0x135dab50>,
<tensorflow.python.keras.layers.core.Dense at 0x135dad90>,
<tensorflow.python.keras.layers.core.Dense at 0x135da580>]
```

In [27]:

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 300)	235500
dense_1 (Dense)	(None, 100)	30100
dense_2 (Dense)	(None, 10)	1010
=====		
Total params: 266,610		
Trainable params: 266,610		

Non-trainable params: 0

In [28]:

```
keras.utils.plot_model(model, "my_fashion_mnist_model.png", show_shapes=True)
```

```
('Failed to import pydot. You must `pip install pydot` and install graphviz  
(https://graphviz.gitlab.io/download/), ', 'for `pydotprint` to work.')
```

In [29]:

```
hidden1 = model.layers[1]  
hidden1.name
```

Out[29]:

```
'dense'
```

In [30]:

```
model.get_layer(hidden1.name) is hidden1
```

Out[30]:

```
True
```

In [31]:

```
weights, biases = hidden1.get_weights()
```

In [32]:

```
weights
```

Out[32]:

```
array([[ 0.02448617, -0.00877795, -0.02189048, ..., -0.02766046,  
        0.03859074, -0.06889391],  
       [ 0.00476504, -0.03105379, -0.0586676 , ...,  0.00602964,  
        -0.02763776, -0.04165364],  
       [-0.06189284, -0.06901957,  0.07102345, ..., -0.04238207,  
        0.07121518, -0.07331658],  
       ...,  
       [-0.03048757,  0.02155137, -0.05400612, ..., -0.00113463,  
        0.00228987,  0.05581069],  
       [ 0.07061854, -0.06960931,  0.07038955, ..., -0.00384101,  
        0.00034875,  0.02878492],  
       [-0.06022581,  0.01577859, -0.02585464, ..., -0.00527829,  
        0.00272203, -0.06793761]], dtype=float32)
```

In [33]:

```
weights.shape
```

Out[33]:

```
(784, 300)
```

In [34]:

```
biases
```

Out[34]:

```
array([0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,  
       0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,  
       0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0., 0.,
```



```
Epoch 1/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.7237 - accuracy: 0.7643 - val_1
oss: 0.5213 - val_accuracy: 0.8226
Epoch 2/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.4842 - accuracy: 0.8317 - val_1
oss: 0.4348 - val_accuracy: 0.8534
Epoch 3/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.4391 - accuracy: 0.8456 - val_1
oss: 0.5347 - val_accuracy: 0.7982
Epoch 4/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.4123 - accuracy: 0.8561 - val_1
oss: 0.3917 - val_accuracy: 0.8646
Epoch 5/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.3940 - accuracy: 0.8620 - val_1
oss: 0.3744 - val_accuracy: 0.8694
Epoch 6/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.3752 - accuracy: 0.8677 - val_1
oss: 0.3707 - val_accuracy: 0.8734
Epoch 7/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.3632 - accuracy: 0.8715 - val_1
oss: 0.3612 - val_accuracy: 0.8722
Epoch 8/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.3517 - accuracy: 0.8747 - val_1
oss: 0.3854 - val_accuracy: 0.8612
Epoch 9/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.3414 - accuracy: 0.8791 - val_1
oss: 0.3586 - val_accuracy: 0.8712
Epoch 10/30
1719/1719 [=====] - 6s 4ms/step - loss: 0.3319 - accuracy: 0.8824 - val_1
```

```
oss: 0.3431 - val_accuracy: 0.8774
Epoch 11/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.3239 - accuracy: 0.8840 - val_l
oss: 0.3449 - val_accuracy: 0.8776
Epoch 12/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.3148 - accuracy: 0.8865 - val_l
oss: 0.3308 - val_accuracy: 0.8820
Epoch 13/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.3080 - accuracy: 0.8894 - val_l
oss: 0.3273 - val_accuracy: 0.8860
Epoch 14/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.3021 - accuracy: 0.8918 - val_l
oss: 0.3414 - val_accuracy: 0.8780
Epoch 15/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.2945 - accuracy: 0.8938 - val_l
oss: 0.3229 - val_accuracy: 0.8838
Epoch 16/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.2890 - accuracy: 0.8969 - val_l
oss: 0.3094 - val_accuracy: 0.8898
Epoch 17/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.2838 - accuracy: 0.8977 - val_l
oss: 0.3539 - val_accuracy: 0.8732
Epoch 18/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.2777 - accuracy: 0.9003 - val_l
oss: 0.3136 - val_accuracy: 0.8896
Epoch 19/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.2727 - accuracy: 0.9021 - val_l
oss: 0.3118 - val_accuracy: 0.8904
Epoch 20/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.2673 - accuracy: 0.9037 - val_l
oss: 0.3281 - val_accuracy: 0.8800
Epoch 21/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.2623 - accuracy: 0.9055 - val_l
oss: 0.3065 - val_accuracy: 0.8916
Epoch 22/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.2576 - accuracy: 0.9076 - val_l
oss: 0.2967 - val_accuracy: 0.8968
Epoch 23/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.2533 - accuracy: 0.9080 - val_l
oss: 0.2983 - val_accuracy: 0.8944
Epoch 24/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.2485 - accuracy: 0.9106 - val_l
oss: 0.3089 - val_accuracy: 0.8902
Epoch 25/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2443 - accuracy: 0.9126 - val_l
oss: 0.2981 - val_accuracy: 0.8940
Epoch 26/30
1719/1719 [=====] - 6s 3ms/step - loss: 0.2405 - accuracy: 0.9137 - val_l
oss: 0.3061 - val_accuracy: 0.8922
Epoch 27/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.2362 - accuracy: 0.9158 - val_l
oss: 0.3010 - val_accuracy: 0.8950
Epoch 28/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.2328 - accuracy: 0.9162 - val_l
oss: 0.2990 - val_accuracy: 0.8948
Epoch 29/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.2284 - accuracy: 0.9186 - val_l
oss: 0.3053 - val_accuracy: 0.8918
Epoch 30/30
1719/1719 [=====] - 5s 3ms/step - loss: 0.2250 - accuracy: 0.9192 - val_l
oss: 0.3027 - val_accuracy: 0.8930
```

In [38]:

```
history.params
```

Out[38]:

```
{'verbose': 1, 'epochs': 30, 'steps': 1719}
```

In [39]:

```
print(history.epoch)
```

```
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29]
```

In [40]:

```
history.history.keys()
```

Out[40]:


```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

In [41]:

```
import pandas as pd

pd.DataFrame(history.history).plot(figsize=(8, 5))
plt.grid(True)
plt.gca().set_ylim(0, 1)
save_fig("keras_learning_curves_plot")
plt.show()
```

Saving figure keras_learning_curves_plot



In [42]:

```
model.evaluate(X_test, y_test)
```

```
313/313 [=====] - 1s 2ms/step - loss: 0.3374 - accuracy: 0.8825
```

Out[42]:

```
[0.33740103244781494, 0.8824999928474426]
```

In [43]:

```
X_new = X_test[:3]
y_proba = model.predict(X_new)
y_proba.round(2)
```

Out[43]:

```
array([[0. , 0. , 0. , 0. , 0. , 0.01, 0. , 0.03, 0. , 0.96],
       [0. , 0. , 0.98, 0. , 0.02, 0. , 0. , 0. , 0. , 0. ],
       [0. , 1. , 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. ]],
      dtype=float32)
```

In [44]:

```
y_pred = model.predict_classes(X_new)
y_pred
```

WARNING:tensorflow:From <ipython-input-44-81ace37e545f>:1: Sequential.predict_classes (from tensorflow.python.keras.engine.sequential) is deprecated and will be removed after 2021-01-01. Instructions for updating:
Please use instead: * `np.argmax(model.predict(x), axis=-1)`, if your model does multi-class classification (e.g. if it uses a `softmax` last-layer activation). * `(model.predict(x) > 0.5).astype("int32")`, if your model does binary classification (e.g. if it uses a `sigmoid` last-layer activation).

Out[44]:

```
array([9, 2, 1], dtype=int64)
```

In [45]:

```
np.array(class_names)[y_pred]
```

Out[45]:

```
array(['Ankle boot', 'Pullover', 'Trouser'], dtype='<U11')
```

In [46]:

```
y_new = y_test[:3]
y_new
```

Out[46]:

```
array([9, 2, 1], dtype=uint8)
```

In [47]:

```
plt.figure(figsize=(7.2, 2.4))
for index, image in enumerate(X_new):
    plt.subplot(1, 3, index + 1)
    plt.imshow(image, cmap="binary", interpolation="nearest")
    plt.axis('off')
    plt.title(class_names[y_test[index]], fontsize=12)
plt.subplots_adjust(wspace=0.2, hspace=0.5)
save_fig('fashion_mnist_images_plot', tight_layout=False)
plt.show()
```

Saving figure fashion_mnist_images_plot



Regresión MLP

Carguemos, dividamos y escalemos el conjunto de datos de viviendas de California:

In [48]:

```
from sklearn.datasets import fetch_california_housing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

housing = fetch_california_housing()
```

```
X_train_full, X_test, y_train_full, y_test = train_test_split(housing.data, housing.target, random_state=42)
X_train, X_valid, y_train, y_valid = train_test_split(X_train_full, y_train_full, random_state=42)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_valid = scaler.transform(X_valid)
X_test = scaler.transform(X_test)
```

In [49]:

```
np.random.seed(42)
tf.random.set_seed(42)
```

In [50]:


```
model = keras.models.Sequential([
    keras.layers.Dense(30, activation="relu", input_shape=X_train.shape[1:]),
    keras.layers.Dense(1)
])
model.compile(loss="mean_squared_error", optimizer=keras.optimizers.SGD(lr=1e-3))
history = model.fit(X_train, y_train, epochs=20, validation_data=(X_valid, y_valid))
mse_test = model.evaluate(X_test, y_test)
X_new = X_test[:3]
y_pred = model.predict(X_new)
```

```
Epoch 1/20
363/363 [=====] - 1s 2ms/step - loss: 1.6419 - val_loss: 0.8560
Epoch 2/20
363/363 [=====] - 1s 1ms/step - loss: 0.7047 - val_loss: 0.6531
Epoch 3/20
363/363 [=====] - 1s 2ms/step - loss: 0.6345 - val_loss: 0.6099
Epoch 4/20
363/363 [=====] - 1s 2ms/step - loss: 0.5977 - val_loss: 0.5658
Epoch 5/20
363/363 [=====] - 0s 1ms/step - loss: 0.5706 - val_loss: 0.5355
Epoch 6/20
363/363 [=====] - 0s 1ms/step - loss: 0.5472 - val_loss: 0.5173
Epoch 7/20
363/363 [=====] - 0s 1ms/step - loss: 0.5288 - val_loss: 0.5081
Epoch 8/20
363/363 [=====] - 0s 1ms/step - loss: 0.5130 - val_loss: 0.4799
Epoch 9/20
363/363 [=====] - 0s 1ms/step - loss: 0.4992 - val_loss: 0.4690
Epoch 10/20
363/363 [=====] - 0s 1ms/step - loss: 0.4875 - val_loss: 0.4656
Epoch 11/20
363/363 [=====] - 0s 1ms/step - loss: 0.4777 - val_loss: 0.4482
Epoch 12/20
363/363 [=====] - 0s 1ms/step - loss: 0.4688 - val_loss: 0.4479
Epoch 13/20
363/363 [=====] - 0s 1ms/step - loss: 0.4615 - val_loss: 0.4296
Epoch 14/20
363/363 [=====] - 0s 1ms/step - loss: 0.4547 - val_loss: 0.4233
Epoch 15/20
363/363 [=====] - 1s 1ms/step - loss: 0.4488 - val_loss: 0.4176
Epoch 16/20
363/363 [=====] - 0s 1ms/step - loss: 0.4435 - val_loss: 0.4123
Epoch 17/20
363/363 [=====] - 0s 1ms/step - loss: 0.4389 - val_loss: 0.4071
Epoch 18/20
363/363 [=====] - 0s 1ms/step - loss: 0.4347 - val_loss: 0.4037
Epoch 19/20
363/363 [=====] - 0s 1ms/step - loss: 0.4306 - val_loss: 0.4000
Epoch 20/20
363/363 [=====] - 0s 1ms/step - loss: 0.4273 - val_loss: 0.3969
162/162 [=====] - 0s 901us/step - loss: 0.4212
```

In [51]:

```
plt.plot(pd.DataFrame(history.history))
plt.grid(True)
plt.gca().set_ylim(0, 1)
```

```
plt.show()
```



```
In [52]:
```

```
y_pred
```

```
Out[52]:
```

```
array([[0.3885664],  
       [1.6792021],  
       [3.1022797]], dtype=float32)
```

```
In [80]:
```

```
class PrintValTrainRatioCallback(keras.callbacks.Callback):  
    def on_epoch_end(self, epoch, logs):  
        print("\nval/train: {:.2f}".format(logs["val_loss"] / logs["loss"]))
```

```
In [81]:
```

```
val_train_ratio_cb = PrintValTrainRatioCallback()  
history = model.fit(X_train, y_train, epochs=1,  
                    validation_data=(X_valid, y_valid),  
                    callbacks=[val_train_ratio_cb])
```

```
360/363 [=====>.] - ETA: 0s - loss: 0.3297
```

```
val/train: 1.08
```

```
363/363 [=====] - 1s 4ms/step - loss: 0.3302 - val_loss: 0.3556
```

```
In [ ]:
```

```
model.evaluate(X_test, y_test)
```